

ML with sklearn

Author: Andrew Sen

Data Exploration

Reading and Manipulating Data

```
In [ ]: import pandas as pd

df = pd.read_csv('auto.csv')
df.head() # print first few rows
```

```
Out[ ]:   mpg  cylinders  displacement  horsepower  weight  acceleration  year  origin  name
0   18.0         8         307.0         130    3504         12.0   70.0     1  chevrolet
chevelle
malibu
1   15.0         8         350.0         165    3693         11.5   70.0     1  buick skylark
320
2   18.0         8         318.0         150    3436         11.0   70.0     1  plymouth
satellite
3   16.0         8         304.0         150    3433         12.0   70.0     1  amc rebel
sst
4   17.0         8         302.0         140    3449         NaN   70.0     1  ford torino
```

```
In [ ]: print('Data frame dimensions: ' + str(df.shape))

Data frame dimensions: (392, 9)
```

```
In [ ]: # print summaries of mpg, weight, year columns
print(df.mpg.describe())
print()
print(df.weight.describe())
print()
print(df.year.describe())
```

```
count    392.000000
mean      23.445918
std        7.805007
min        9.000000
25%       17.000000
50%       22.750000
75%       29.000000
max       46.600000
Name: mpg, dtype: float64
```

```
count    392.000000
mean    2977.584184
std     849.402560
min    1613.000000
25%    2225.250000
50%    2803.500000
75%    3614.750000
max    5140.000000
Name: weight, dtype: float64
```

```
count    390.000000
mean      76.010256
std        3.668093
min       70.000000
25%       73.000000
50%       76.000000
75%       79.000000
max       82.000000
Name: year, dtype: float64
```

```
||Range|Average ---|---|--- mpg|37.6|23.445918 weight|3527|2977.584184 year|12|76.010256
```

```
In [ ]: df.dtypes # print types of columns
```

```
Out[ ]: mpg          float64
cylinders          int64
displacement       float64
horsepower         int64
weight             int64
acceleration       float64
year               float64
origin             int64
name               object
dtype: object
```

```
In [ ]: df.cylinders = df.cylinders.astype('category').cat.codes
df.origin = df.origin.astype('category')
df.dtypes
```

```
Out[ ]: mpg          float64
cylinders         int8
displacement      float64
horsepower        int64
weight            int64
acceleration      float64
year              float64
origin            category
name              object
dtype: object
```

```
In [ ]: df.dropna(inplace=True) # drop rows with NA values
print('Dimensions after dropping NAs: ' + str(df.shape))
```

Dimensions after dropping NAs: (389, 9)

```
In [ ]: # make binary classifier 'mpg_high' column
avg_mpg = df.mpg.mean()
df['mpg_high'] = df.mpg.map(lambda x: 1 if x > avg_mpg else 0).astype('category')

# remove 'mpg' and 'name' columns
df.drop(columns=['mpg', 'name'], inplace=True)

df.head()
```


```
Out[ ]:   cylinders  displacement  horsepower  weight  acceleration  year  origin  mpg_high
0         4          307.0         130    3504         12.0  70.0      1         0
1         4          350.0         165    3693         11.5  70.0      1         0
2         4          318.0         150    3436         11.0  70.0      1         0
3         4          304.0         150    3433         12.0  70.0      1         0
6         4          454.0         220    4354          9.0  70.0      1         0
```

Graphs

```
In [ ]: import seaborn as sb

sb.catplot(x='mpg_high', kind='count', data=df)
```


```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x2aa260aacd0>
```



This graph shows that most cars in the dataset have below average mpg.

```
In [ ]: sb.relplot(x='horsepower', y='weight', hue='mpg_high', data=df)
```


```
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x2aa25f98a90>
```



This graph shows that vehicles with high horsepower and weight tend to have below average mpg.

```
In [ ]: sb.boxplot(x='mpg_high', y='weight', data=df)
```

```
Out[ ]: <AxesSubplot: xlabel='mpg_high', ylabel='weight'>
```



This graph shows that cars with above average mpg tend to have lower weight than cars with below average mpg.

Classification

We will now perform classification using various methods available in sklearn. First, we will divide the data into train/test data.

```
In [ ]: from sklearn.model_selection import train_test_split

X = df.loc[:, df.columns != 'mpg_high']
y = df.mpg_high
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print('Dimensions:')
print(f'X_train: {str(X_train.shape)}')
print(f'y_train: {str(y_train.shape)}')
print(f'X_test: {str(X_test.shape)}')
print(f'y_test: {str(y_test.shape)}')
```

Dimensions:
X_train: (311, 7)
y_train: (311,)
X_test: (78, 7)
y_test: (78,)

Logistic Regression

```
In [ ]: from sklearn.metrics import classification_report
        from sklearn.linear_model import LogisticRegression
```

```
lr = LogisticRegression(solver='lbfgs')
lr.fit(X_train, y_train)
pred_lr = lr.predict(X_test)
print(classification_report(y_test, pred_lr))
```

	precision	recall	f1-score	support
0	0.98	0.80	0.88	50
1	0.73	0.96	0.83	28
accuracy			0.86	78
macro avg	0.85	0.88	0.85	78
weighted avg	0.89	0.86	0.86	78

C:\Users\andre\AppData\Roaming\Python\Python311\site-packages\sklearn\linear_model_logistic.py:444: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
In [ ]:
```

Accuracy is 86%. Precision is higher for class 0, but the recall is higher for class 1. We would like both precision and recall to be more even between both classes. However, the training function hit the maximum number of iterations, so the algorithm may have been able to perform better in ideal conditions.

Decision Trees

```
In [ ]: from sklearn.tree import DecisionTreeClassifier
```

```
dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
pred_dt = dt.predict(X_test)
print(classification_report(y_test, pred_dt))
```

```
# plotting tree
```

```
from sklearn import tree
tree.plot_tree(dt)
```


	precision	recall	f1-score	support
0	0.96	0.90	0.93	50
1	0.84	0.93	0.88	28
accuracy			0.91	78
macro avg	0.90	0.91	0.90	78
weighted avg	0.91	0.91	0.91	78


```
Out[ ]: [Text(0.6433823529411765, 0.9444444444444444, 'X[0] <= 2.5\ngini = 0.5\nsamples = 3
11\nvalue = [153, 158]'),
Text(0.4338235294117647, 0.8333333333333334, 'X[2] <= 101.0\ngini = 0.239\nsamples
= 173\nvalue = [24, 149]'),
Text(0.27941176470588236, 0.7222222222222222, 'X[5] <= 75.5\ngini = 0.179\nsamples
= 161\nvalue = [16, 145]'),
Text(0.14705882352941177, 0.6111111111111112, 'X[1] <= 119.5\ngini = 0.362\nsample
s = 59\nvalue = [14, 45]'),
Text(0.058823529411764705, 0.5, 'X[4] <= 13.75\ngini = 0.159\nsamples = 46\nvalue
= [4, 42]'),
Text(0.029411764705882353, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue =
[2, 0]'),
Text(0.08823529411764706, 0.3888888888888889, 'X[3] <= 2683.0\ngini = 0.087\nsampl
es = 44\nvalue = [2, 42]'),
Text(0.058823529411764705, 0.2777777777777778, 'X[3] <= 2377.0\ngini = 0.045\nsamp
les = 43\nvalue = [1, 42]'),
Text(0.029411764705882353, 0.1666666666666666, 'gini = 0.0\nsamples = 38\nvalue =
[0, 38]'),
Text(0.08823529411764706, 0.1666666666666666, 'X[3] <= 2385.0\ngini = 0.32\nsampl
es = 5\nvalue = [1, 4]'),
Text(0.058823529411764705, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue =
[1, 0]'),
Text(0.11764705882352941, 0.05555555555555555, 'gini = 0.0\nsamples = 4\nvalue =
[0, 4]'),
Text(0.11764705882352941, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue =
[1, 0]'),
Text(0.23529411764705882, 0.5, 'X[4] <= 17.75\ngini = 0.355\nsamples = 13\nvalue =
[10, 3]'),
Text(0.20588235294117646, 0.3888888888888889, 'X[2] <= 81.5\ngini = 0.469\nsamples
= 8\nvalue = [5, 3]'),
Text(0.17647058823529413, 0.2777777777777778, 'gini = 0.0\nsamples = 2\nvalue =
[0, 2]'),
Text(0.23529411764705882, 0.2777777777777778, 'X[5] <= 71.5\ngini = 0.278\nsamples
= 6\nvalue = [5, 1]'),
Text(0.20588235294117646, 0.1666666666666666, 'X[2] <= 88.0\ngini = 0.5\nsamples
= 2\nvalue = [1, 1]'),
Text(0.17647058823529413, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue =
[1, 0]'),
Text(0.23529411764705882, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue =
[0, 1]'),
Text(0.2647058823529412, 0.1666666666666666, 'gini = 0.0\nsamples = 4\nvalue =
[4, 0]'),
Text(0.2647058823529412, 0.3888888888888889, 'gini = 0.0\nsamples = 5\nvalue = [5,
0]'),
Text(0.4117647058823529, 0.6111111111111112, 'X[3] <= 3250.0\ngini = 0.038\nsampl
es = 102\nvalue = [2, 100]'),
Text(0.35294117647058826, 0.5, 'X[3] <= 2880.0\ngini = 0.02\nsamples = 100\nvalue
= [1, 99]'),
Text(0.3235294117647059, 0.3888888888888889, 'gini = 0.0\nsamples = 94\nvalue =
[0, 94]'),
Text(0.38235294117647056, 0.3888888888888889, 'X[3] <= 2920.0\ngini = 0.278\nsampl
es = 6\nvalue = [1, 5]'),
Text(0.35294117647058826, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue =
[1, 0]'),
Text(0.4117647058823529, 0.2777777777777778, 'gini = 0.0\nsamples = 5\nvalue = [0,
5]'),
```

```

Text(0.47058823529411764, 0.5, 'X[3] <= 3400.0\ngini = 0.5\nsamples = 2\nvalue =
[1, 1]'),
Text(0.4411764705882353, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [1,
0]'),
Text(0.5, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.5882352941176471, 0.7222222222222222, 'X[4] <= 14.45\ngini = 0.444\nsamples
= 12\nvalue = [8, 4]'),
Text(0.5588235294117647, 0.6111111111111112, 'X[5] <= 76.0\ngini = 0.444\nsamples
= 6\nvalue = [2, 4]'),
Text(0.5294117647058824, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(0.5882352941176471, 0.5, 'X[4] <= 13.95\ngini = 0.444\nsamples = 3\nvalue =
[2, 1]'),
Text(0.5588235294117647, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [2,
0]'),
Text(0.6176470588235294, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [0,
1]'),
Text(0.6176470588235294, 0.6111111111111112, 'gini = 0.0\nsamples = 6\nvalue = [6,
0]'),
Text(0.8529411764705882, 0.8333333333333334, 'X[5] <= 79.5\ngini = 0.122\nsamples
= 138\nvalue = [129, 9]'),
Text(0.7941176470588235, 0.7222222222222222, 'X[4] <= 21.6\ngini = 0.045\nsamples
= 129\nvalue = [126, 3]'),
Text(0.7647058823529411, 0.6111111111111112, 'X[3] <= 2737.0\ngini = 0.031\nsample
s = 128\nvalue = [126, 2]'),
Text(0.7058823529411765, 0.5, 'X[4] <= 13.45\ngini = 0.444\nsamples = 3\nvalue =
[2, 1]'),
Text(0.6764705882352942, 0.3888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [0,
1]'),
Text(0.7352941176470589, 0.3888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [2,
0]'),
Text(0.8235294117647058, 0.5, 'X[2] <= 83.0\ngini = 0.016\nsamples = 125\nvalue =
[124, 1]'),
Text(0.7941176470588235, 0.3888888888888889, 'X[1] <= 225.0\ngini = 0.375\nsamples
= 4\nvalue = [3, 1]'),
Text(0.7647058823529411, 0.2777777777777778, 'gini = 0.0\nsamples = 1\nvalue = [0,
1]'),
Text(0.8235294117647058, 0.2777777777777778, 'gini = 0.0\nsamples = 3\nvalue = [3,
0]'),
Text(0.8529411764705882, 0.3888888888888889, 'gini = 0.0\nsamples = 121\nvalue =
[121, 0]'),
Text(0.8235294117647058, 0.6111111111111112, 'gini = 0.0\nsamples = 1\nvalue = [0,
1]'),
Text(0.9117647058823529, 0.7222222222222222, 'X[1] <= 196.5\ngini = 0.444\nsamples
= 9\nvalue = [3, 6]'),
Text(0.8823529411764706, 0.6111111111111112, 'gini = 0.0\nsamples = 4\nvalue = [0,
4]'),
Text(0.9411764705882353, 0.6111111111111112, 'X[1] <= 247.0\ngini = 0.48\nsamples
= 5\nvalue = [3, 2]'),
Text(0.9117647058823529, 0.5, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(0.9705882352941176, 0.5, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]')

```



This is slightly more accurate than the logistic regression model at 88%. The values for precision and recall are also more even between both classes.

Neural Network

Before making a neural network, we'll first scale the data to improve the performance of our models.

```
In [ ]: from sklearn import preprocessing

scaler = preprocessing.StandardScaler().fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Now we'll make our first neural network. The choice of network topology is mostly arbitrary at this stage.

```
In [ ]: from sklearn.neural_network import MLPClassifier

nn1 = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(6, 3), max_iter=1500, random_state=1)
nn1.fit(X_train_scaled, y_train)
pred_nn1 = nn1.predict(X_test_scaled)
print(classification_report(y_test, pred_nn1))
```

	precision	recall	f1-score	support
0	0.96	0.88	0.92	50
1	0.81	0.93	0.87	28
accuracy			0.90	78
macro avg	0.88	0.90	0.89	78
weighted avg	0.90	0.90	0.90	78

With these settings, we'll already gotten 90% accuracy, outperforming the other two models, and precision and recall are fairly balanced.

Now we will make a second neural network. For this model, I will change the solver and increase the number of layers to see if it will get better results.

```
In [ ]: nn2 = MLPClassifier(solver='sgd', hidden_layer_sizes=(7, 5, 2), max_iter=1500, random_state=42)
nn2.fit(X_train_scaled, y_train)
pred_nn2 = nn2.predict(X_test_scaled)
print(classification_report(y_test, pred_nn2))
```

	precision	recall	f1-score	support
0	0.98	0.82	0.89	50
1	0.75	0.96	0.84	28
accuracy			0.87	78
macro avg	0.86	0.89	0.87	78
weighted avg	0.89	0.87	0.87	78

As we can see, this model unfortunately performed worse than the first neural network. This may be because the addition of a third hidden layer caused the model to overfit the training data. Since the dataset is relatively small, overfitting is very likely.

Analysis

The decision tree performed the best, followed by the first neural network, the second neural network, and finally the logistic regression model.

With regards to class 0, the logistic regression model and second neural network had the highest precision, followed by the decision tree and first neural network. The decision tree had the highest recall, followed by the first neural network, second neural network, and finally the logistic regression model.

With regards to class 1, the decision tree had the highest precision, followed by the first neural network, second neural network, and finally the logistic regression model. The logistic regression model and second neural network had the highest recall, followed by the decision tree and first neural network. The decision tree overall had the most balanced values for both precision and recall.

Both neural networks may have slightly overfit the data given that the training set is relatively small. The logistic regression algorithm also failed to converge, which may be the reason it underperformed compared to decision trees. However, the decision tree may have been able to better capture the data regardless.

Overall, I much prefer the experience of using pandas and sklearn in Python compared to writing R. Pandas feels easier to learn compared to R's syntax for data manipulation. Since I am writing this in an IDE, I get tooltips for text-completion and documentation. While RStudio has text completion, learning how to use its various functionalities usually necessitates going to the help console or simply looking things up online.